


Explanation techniques for neural networks

Laura Rieger (lauri@dtu.dk)



Colab notebook

- Colab notebook:
 - **bit.ly/2PdScxx**
 - File -> Save a Copy in Drive
 - Edit -> Notebook settings -> GPU
 - Tools -> Preferences -> Misc ->

Notebook settings

Runtime type
Python 3

Hardware accelerator
GPU

☒ Omit code cell output when saving this notebook

CANCEL SAVE

Power Level
Some power

☐ Corgi Mode


☐ Kitty Mode

Intro

- Assumptions
 - Complicated non-linear task
 - Can **not** be solved by an intuitively explainable model
- Complete understanding is not possible
 - Single decision explanations
 - Region-based explanations
- Approaches
 - Backpropagation
 - Local approximation with simple model

Neural Networks

- Complex function composed of smaller functions
- Trained by large amounts of labelled data
- Training:
 - Differentiate loss over weights
 - Update weights
- Differentiation done automatically



Input Layer $\in \mathbb{R}^3$

Hidden Layer $\in \mathbb{R}^5$

Output Layer $\in \mathbb{R}^2$

$$\mathcal{L}(\theta, X, y) = f_{\theta}(X) * \log(y) + (1 - f_{\theta}(X)) * \log(1 - y)$$


$$\Delta\theta = \lambda * \frac{\delta\mathcal{L}(\theta, X, y)}{\theta}$$

Saliency [1]

- Want to know what parts of the input responsible for the output
- Areas with a high gradient
 - Output changes fast here
 - Probably important
- Output is fully differentiable
- Basic explainability method
- Super noisy
 - High redundancy
 - Noisy function
- Popular variant:
Guided Backprop [9]

$$\frac{\delta y_c}{\delta X}$$

$$\frac{\partial' y_c}{\partial' x_{i,j}} = \text{ReLU} \left(\frac{\partial y_c}{\partial l_{-1}} \right) \dots \text{ReLU} \left(\frac{\partial l_1}{\partial x_{i,j}} \right)$$




SmoothGrad [15]

- How do we remove noise? With more noise!

$$\sum_N \frac{\delta y_c}{\delta(X + \sigma_n)}$$

- Average over number of input with noise added
- Clears up image
- Applicable to all basic explanation techniques




Integrated Gradients [16]

- How do we remove noise? With more (different) noise!


$$\frac{1}{N} \sum_N \frac{\delta y_c}{\delta(\mathbf{0} + o_n * (X - \mathbf{0}))} * (X - \mathbf{0}), o_n \sim U(0, 1)$$

- Average over samples on the pathway between baseline
- Dependent on choice of baseline (all black, average, ...)
- Recent extension: use samples as baseline



GradCAM [13]

- Gradient-weighted Class Activation Mapping
- Way less noisy
- No finegrained input
- Procedure:
 - Calculate gradient to last convolutional layer
 - Average gradient for each channel
 - Multiply average gradient with all activations
 - Average over all channels
 - Upsample




Other methods based on backpropagation

- LRP [12]
- Expected Gradients [17]
- CAM [2]
- GuidedBackprop [9]

Local approximation with interpretable model – LIME [4]

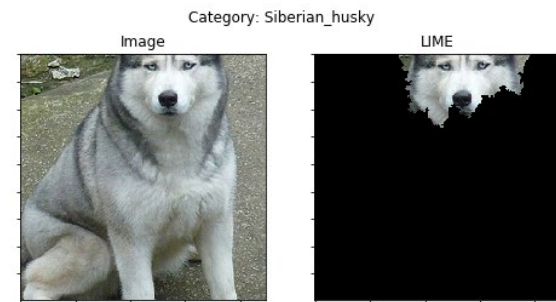

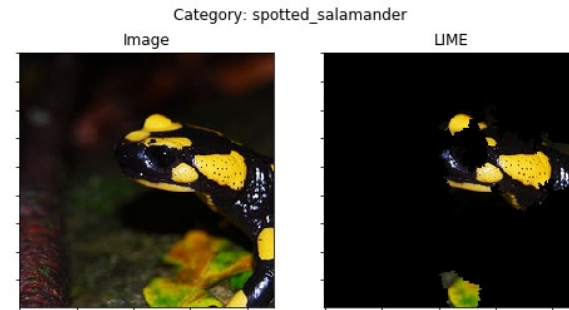
- Intuition:
 - Sample around \mathbf{x}
 - Weigh samples according to distance
 - Train linear classifier
 - Obtain explanation
- Low-dimensional representation necessary
 - For images: segment into super-pixels
 - For text: bag of words



From <https://github.com/marcotcr/lime>

Local approximation with interpretable model – LIME [4]

- Intuition:
 - Sample around \mathbf{x}
 - Weigh samples according to distance
 - Train linear classifier
 - Obtain explanation
- Low-dimensional representation necessary
 - For images: segment into super-pixels
 - For text: bag of words



Images obtained with LIME library
from pretrained VGG16


Higher-level

- Network level explanations
- Requires domain knowledge
- Interesting for risk and fairness analysis
- Two approaches presented
 - Analyzing specific network parts
 - Analyzing specific aspects

Probing the network

- “Understanding Neural Networks Through Deep Visualization” [5]
 - Idea: iteratively optimize activation of neurons with backpropagation
 - Regularize to encourage realism
 - For output or intermediate layers
- Alternatives
 - Bau, David, et al. "Network Dissection: Quantifying Interpretability of Deep Visual Representations." *Computer Vision and Pattern Recognition (CVPR), 2017 IEEE Conference on*. IEEE, 2017.
 - Alain, Guillaume, and Yoshua Bengio. "Understanding intermediate layers using linear classifier probes." (2016).

Category: hen



Obtained with kerasvis
from pretrained VGG16

Testing with Concept Activation Vectors CAV

- Proposed by Kim et al [8]
- Requires high domain knowledge
- Idea:
 - assemble dataset $\{P, N\}$
 - Train linear classifier on representation of given layer
 - Obtain weights
- Useful for
 - Evaluating given input
 - Identifying a known bias in dataset and model

Model Women concept: most similar necktie images



Model Women concept: least similar necktie images




Kim, Been, et al.

"TCAV: Relative concept importance testing with Linear Concept Activation Vectors." (2018).

What should I use?

- How do we evaluate interpretation methods?
 - Humans are really bad evaluators
- Theoretical properties
- Sanity checks
- Human evaluation
 - Fooled by clearness
- Remove input consecutively – check output degradation



Integrated Gradients on random weights

Take-away

- We have to make a trade-off when obtaining explanations from NNs
- Different approaches have different pros and contras
- Task in question needs to be considered

	Fidelity	Understandability	Sufficiency	Low construction overhead	Efficiency
Backprop	+	-	0	+	+
Local	0	+	0	+	-
High-level	+	+	-	-	0

References

1. Simonyan, Karen, Andrea Vedaldi, and Andrew Zisserman. "Deep inside convolutional networks: Visualising image classification models and saliency maps." *arXiv preprint arXiv:1312.6034* (2013).
2. Zhou, Bolei, et al. "Learning deep features for discriminative localization." *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*. 2016.
3. Zhang, Quanshi, Ying Nian Wu, and Song-Chun Zhu. "Interpretable Convolutional Neural Networks."
4. Ribeiro, Marco Tulio, Sameer Singh, and Carlos Guestrin. "Why should i trust you?: Explaining the predictions of any classifier." *Proceedings of the 22nd ACM SIGKDD international conference on knowledge discovery and data mining*. ACM, 2016.
5. Yosinski, Jason, et al. "Understanding neural networks through deep visualization." In *ICML Workshop on Deep Learning*.
6. Bau, David, et al. "Network dissection: Quantifying interpretability of deep visual representations." *arXiv preprint arXiv:1704.05796* (2017).
7. Alain, Guillaume, and Yoshua Bengio. "Understanding intermediate layers using linear classifier probes." *arXiv preprint arXiv:1610.01644* (2016).
8. Kim, Been, et al. "TCAV: Relative concept importance testing with Linear Concept Activation Vectors." (2018).
9. Springenberg, Jost Tobias, et al. "Striving for simplicity: The all convolutional net." *arXiv preprint arXiv:1412.6806* (2014).
10. Zhang Q, Wu YN, Zhu S-C. Interpretable Convolutional Neural Networks. <https://arxiv.org/pdf/1710.00935.pdf>. Accessed February 20, 2018.
11. Smilkov, Daniel, et al. "Smoothgrad: removing noise by adding noise." *arXiv preprint arXiv:1706.03825* (2017).
12. Bach, Sebastian, et al. "On pixel-wise explanations for non-linear classifier decisions by layer-wise relevance propagation." *PloS one* 10.7 (2015): e0130140.
13. Selvaraju, Ramprasaath R., et al. "Grad-CAM: Visual Explanations from Deep Networks via Gradient-Based Localization." *ICCV*. 2017.
14. Swartout, William R., and Johanna D. Moore. "Explanation in second generation expert systems." *Second generation expert systems*. Springer, Berlin, Heidelberg, 1993. 543-585.
15. Smilkov, Daniel, et al. "Smoothgrad: removing noise by adding noise." *arXiv preprint arXiv:1706.03825* (2017).
16. Sundararajan, Mukund, Ankur Taly, and Qiqi Yan. "Axiomatic attribution for deep networks." *Proceedings of the 34th International Conference on Machine Learning-Volume 70*. JMLR. org, 2017.
17. Erion, Gabriel, et al. "Learning Explainable Models Using Attribution Priors." *arXiv preprint arXiv:1906.10670* (2019).